

*# This jupyter notebook is prepared by Nawras Rawas Qalaji*

```
import pandas as pd
import numpy as np
import matplotlib as mpl
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
import missingno as msno
import scipy.stats as st
from sklearn import ensemble, tree, linear_model
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, PolynomialFeatures
from sklearn.linear_model import LinearRegression, SGDRegressor,
Ridge, ElasticNet, Lasso
from sklearn.metrics import mean_squared_error, mean_absolute_error
from sklearn.metrics import mean_squared_error
from sklearn.pipeline import Pipeline
```

*#import hrdata.csv to hrdata dataframe*

```
hrdata = pd.read_csv("/content/drive/MyDrive/Colab
Notebooks/ecommarce.csv")
```

*#count columns and rows*

```
print("Count Columns and rows")
print("columns " + str(hrdata.shape[1]))
print("rows " + str(hrdata.shape[0]) + "\n")
```

*#describe datafram*

```
print("\nDescribe")
print(hrdata.describe())
```

```
Count Columns and rows
columns 9
rows 500
```

Describe

	Unnamed: 0	Credit Card	Avg. Session Length	Time on App	\
count	500.000000	5.000000e+02	500.000000	500.000000	
mean	249.500000	3.706324e+17	34.053194	13.052488	
std	144.481833	1.235588e+18	0.992563	0.994216	
min	0.000000	5.018057e+11	30.532429	9.508152	
25%	124.750000	3.683275e+13	33.341822	12.388153	
50%	249.500000	3.513612e+15	34.082008	12.983231	
75%	374.250000	4.777131e+15	34.711985	13.753850	
max	499.000000	4.959148e+18	37.139662	16.126994	

Time on Website   Length of Membership   Yearly Amount Spent

count	500.000000	500.000000	500.000000
mean	38.060445	4.033462	500.314038
std	1.010489	0.999278	79.314782
min	34.913847	0.769901	257.670582
25%	37.349257	3.430450	446.038277
50%	38.069367	4.033975	499.887875
75%	38.716432	4.626502	550.313828
max	41.005182	7.422689	766.518462

Explain in words about the description of any two variables

- The data collected on the credit card is not really useable since its just the credit card number instead of any actual data related to it
- The data on unnamed is also useless since its just the index of the array

*#top and bottom 5 rows*

```
print("\nTop Five Rows")
print(hrdata.loc[[0,1,2,3,4]])
print("\nBottom Five Rows")
print(hrdata.iloc[[-1,-2,-3,-4,-5]])
```

*#numerical columns*

```
print("\nNumerical Columns")
print(hrdata.select_dtypes(include = [np.number]).columns)
```

*#categorical columns*

```
print("\nCategorical Columns")
print(hrdata.select_dtypes(include = [np.object]).columns)
```

*#missing values*

```
print("\nMissing Values Numerically")
missingValues = hrdata.isnull().sum().sort_values(ascending = False)
print(missingValues)
```

```
print("\nMissing Values Percentage")
print(missingValues/len(hrdata)*100)
```

Top Five Rows

Unnamed: 0	Email	Address \
0	adkv@ota.com	89280 Mark Lane\nNew John, MN 16131
1	gjun@syj.com	363 Amanda Cliff Apt. 638\nWest Angela, KS 31437
2	qjyr@pkk.com	62008 Adam Lodge\nLake Pamela, NY 30677
3	jkiu@xsb.com	950 Tami Island\nLake Aimeeview, MT 93614
4	stvb@niy.com	08254 Kelly Squares\nNorth Lauren,

AR 78382

Website \	Credit Card	Avg. Session Length	Time on App	Time on
0	3544288738428794	35.497268	13.655651	
40.577668				
1	6546228325389133	32.926272	12.109461	
38.268959				
2	4406395951712628314	34.000915	12.330278	
38.110597				
3	30334036663133	35.305557	14.717514	
37.721283				
4	3582080469154498	34.330673	13.795189	
38.536653				

	Length of Membership	Yearly Amount Spent
0	4.582621	588.951054
1	3.164034	393.204933
2	4.604543	488.547505
3	3.620179	582.852344
4	4.946308	600.406092

Bottom Five Rows

Unnamed: 0	Email \
499	499 phqb@nlg.com
498	498 zvtz@onj.com
497	497 pndt@jyr.com
496	496 awrc@iok.com
495	495 xskz@gwj.com

Card \	Address	Credit
499	424 Mark Junctions\nDarrellchester, TX 09088	
5427200269739116		
498	5568 Robert Station Apt. 030\nTurnerstad, GA 9...	
36218092488069		
497	1555 Chen Road\nBergerchester, NH 46418	
4086276267550896697		
496	663 Christopher Garden\nLake Carrieberg, PA 70796	
6011536844623717		
495	7083 Wallace Rest\nNew Trevor, NM 70240	
30206742023085		

Membership \	Avg. Session Length	Time on App	Time on Website	Length of
499	34.715981	13.418808	36.771016	
3.235160				
498	34.322501	13.391423	37.840086	
2.836485				
497	33.646777	12.499409	39.332576	

5.458264			
496	35.702529	12.695736	38.190268
4.076526			
495	34.237660	14.566160	37.417985
4.246573			

Yearly Amount Spent	
499	498.778642
498	457.469510
497	552.620145
496	530.049004
495	574.847438

#### Numerical Columns

```
Index(['Unnamed: 0', 'Credit Card', 'Avg. Session Length', 'Time on App',
      'Time on Website', 'Length of Membership', 'Yearly Amount Spent'],
      dtype='object')
```

#### Categorical Columns

```
Index(['Email', 'Address'], dtype='object')
```

#### Missing Values Numerically

Unnamed: 0	0
Email	0
Address	0
Credit Card	0
Avg. Session Length	0
Time on App	0
Time on Website	0
Length of Membership	0
Yearly Amount Spent	0

dtype: int64

#### Missing Values Percentage

Unnamed: 0	0.0
Email	0.0
Address	0.0
Credit Card	0.0
Avg. Session Length	0.0
Time on App	0.0
Time on Website	0.0
Length of Membership	0.0
Yearly Amount Spent	0.0

dtype: float64

```
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:13:
DeprecationWarning: `np.object` is a deprecated alias for the builtin
`object`. To silence this warning, use `object` by itself. Doing this
```


will not modify any behavior and is safe.  
Deprecated in NumPy 1.20; for more details and guidance:  
<https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations>  
`del sys.path[0]`

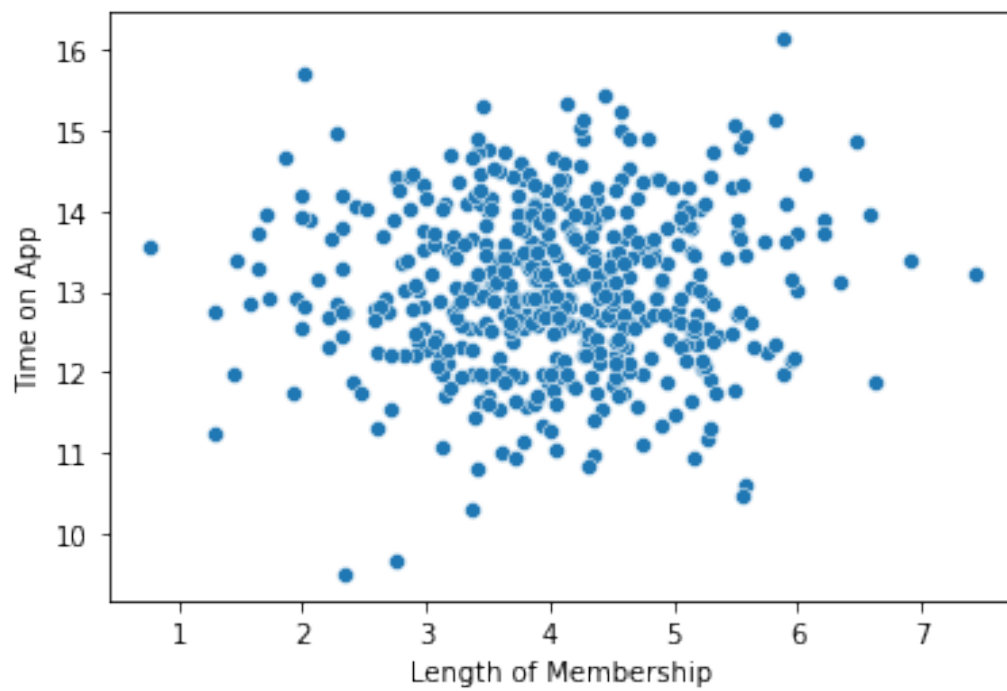

*#scatterplots*

```
plt.figure()  
sns.scatterplot(x = "Yearly Amount Spent", y = "Time on Website", data  
= hrdata)  
plt.show()
```

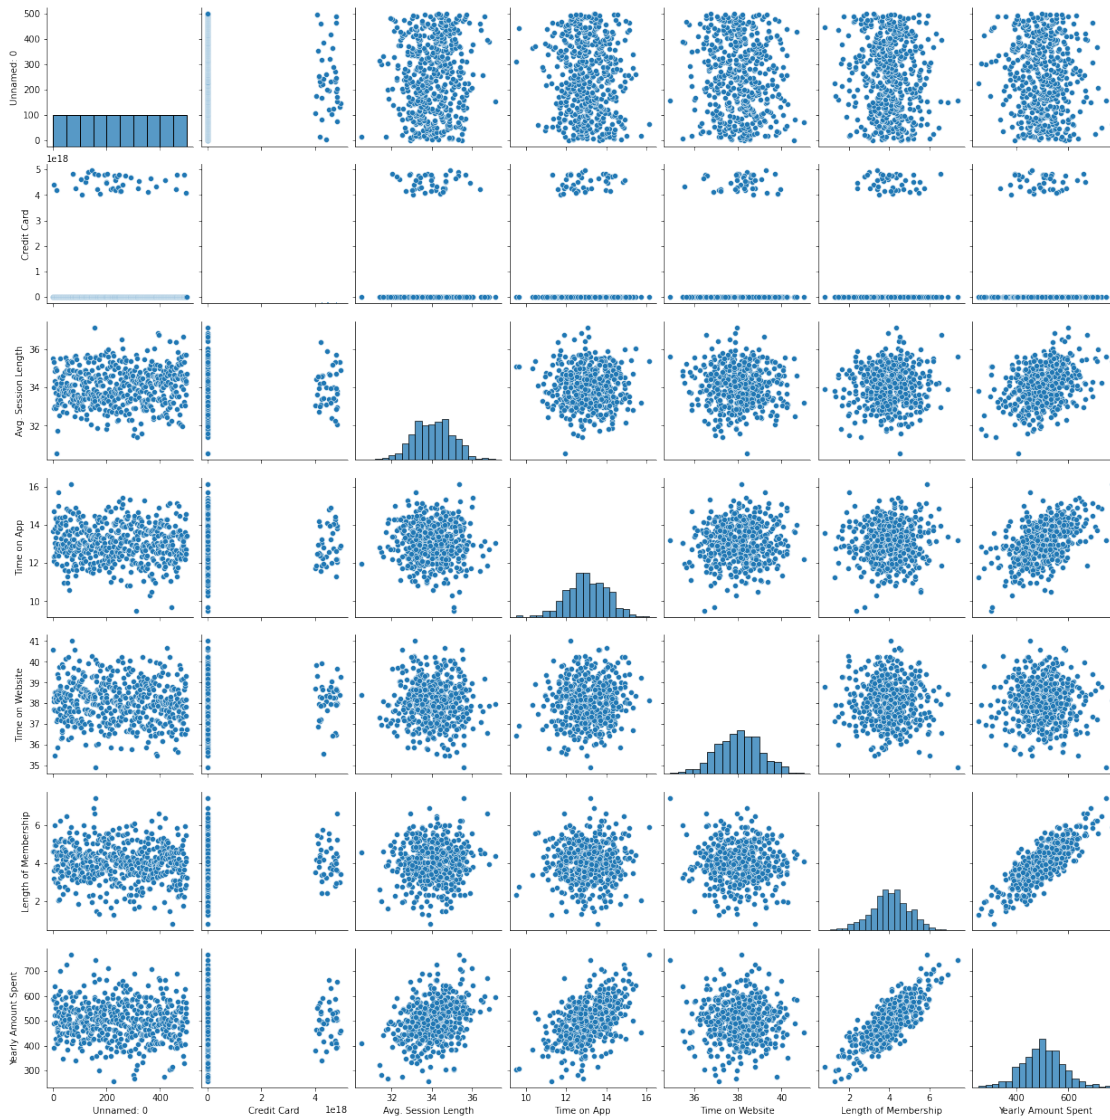
```
plt.figure()  
sns.scatterplot(x = "Yearly Amount Spent", y = "Time on App", data =  
hrdata)  
plt.show()
```

```
plt.figure()  
sns.scatterplot(x = "Length of Membership", y = "Time on App", data =  
hrdata)  
plt.show()
```





```
sns.pairplot(data = hrdata)
plt.figure()
plt.show()
```



<Figure size 432x288 with 0 Axes>

```
correlationNums = hrdata.select_dtypes(include = [np.number])
sns.heatmap(data = correlationNums.corr(), annot = True)
plt.figure()
plt.show()
```




<Figure size 432x288 with 0 Axes>

discuss which columns must be removed based on that and which column is mostly interesting and related to Yearly Amount Spent?

- You would probably drop Unnamed: 0, credit card, and time on website because they have a negative correlation. Length of membership is most interesting and related to yearly amount spent.

```
plt.figure()
sns.scatterplot(x = "Yearly Amount Spent", y = "Length of Membership",
data = hrdata)
plt.show()
```





*#drop unnecessary columns*

```
hrdata.drop(["Unnamed: 0", "Credit Card", "Email", "Address"], axis = 1)
```

	Avg. Session Length	Time on App	Time on Website	Length of Membership \
0	35.497268	13.655651	40.577668	4.582621
1	32.926272	12.109461	38.268959	3.164034
2	34.000915	12.330278	38.110597	4.604543
3	35.305557	14.717514	37.721283	3.620179
4	34.330673	13.795189	38.536653	4.946308
..	...	...	...	...
495	34.237660	14.566160	37.417985	4.246573
496	35.702529	12.695736	38.190268	4.076526
497	33.646777	12.499409	39.332576	5.458264
498	34.322501	13.391423	37.840086	2.836485
499	34.715981	13.418808	36.771016	3.235160

	Yearly Amount Spent
0	588.951054
1	393.204933
2	488.547505
3	582.852344
4	600.406092
...	...
495	574.847438
496	530.049004
497	552.620145
498	457.469510
499	498.778642

[500 rows x 5 columns]

```
X = hrdata[["Avg. Session Length", "Time on App", "Time on Website",
"Length of Membership"]]
y = hrdata["Yearly Amount Spent"]
```

```
scaler = StandardScaler()
scaler.fit(X)
X = scaler.transform(X)
```

```
X_train, X_test, y_train, y_test = train_test_split(X,y, test_size =
.3, random_state = 101)
```

```
#Linear regression
lr = LinearRegression()
```

```
#train
lr.fit(X_train, y_train)
```

```
#predict
y_pred = lr.predict(X_test)
```

```
#print coeffecients and intercepts
print("Coefficients " + str(lr.coef_))
print("Intercept " + str(lr.intercept_))
```

```
#generate scatterplot
plt.xlabel("y_test")
plt.ylabel("y_pred", rotation = 0)
plt.scatter(y_test, y_pred)
plt.show()
```

```
# Compute and print R^2, RMSE, MAE and MSE
print("R^2: {}".format(lr.score(X_test, y_test)))
```

```
print("Mean Squared Error: {}".format(mean_squared_error(y_test,
```

```


y_pred)))

print("Mean Absolute Error: {}".format(mean_absolute_error(y_test,
y_pred)))

rmse = np.sqrt(mean_squared_error(y_test, y_pred))
print("Root Mean Squared Error: {}".format(rmse))

Coefficients [25.76252659 38.32855202  0.19220992 61.17355707]
Intercept 500.5316677457304

```



```

R^2: 0.9890046246741234
Mean Squared Error: 79.8130516509746
Mean Absolute Error: 7.228148653430835
Root Mean Squared Error: 8.933815066978642

```

Interpret the coefficient and which coefficient belongs to which feature and based on that explain any strategy that should help the business

- 25.7625 belongs to Avg session length, 38.3285 belongs to time on app, 0.1922 belongs to time on website and 61.1735 belongs to length of membership.
- A higher coefficient means a stronger correlation so its expected that those with a longer membership will have a larger yearly amount spent and the business should focus on those with a longer membership.

*#Normal Equation*

*#add 1s*

```
X_b = np.c_[np.ones((350,1)), X_train]
```

```
#change y to numpy and reshape
```

```
y_new = y_train.to_numpy()
```

```
y_new = y_new.reshape(-1,1)
```

```
#normal equation
```

```
theta_best_values=np.linalg.inv(X_b.T.dot(X_b)).dot(X_b.T).dot(y_new)
```

```
print("theta values\n" + str(theta_best_values))
```

```
theta values
```

```
[[5.00531668e+02]
```

```
 [2.57625266e+01]
```

```
 [3.83285520e+01]
```

```
 [1.92209922e-01]
```

```
 [6.11735571e+01]]
```

Are they very close to the sklearn's linear regression

- Theta best values are very close to the intercept and coefficients of the linear regression

LR Intercept 500.53620759592945

LR Coefficients [25.76851649 38.34767562 61.16352344]

Theta [[500.5362076 ] [ 25.76851649] [ 38.34767562] [ 61.16352344]]

```
#Normal equation Linear Regression
```

```
X_test_new = np.c_[np.ones((len(X_test),1)), X_test]
```

```
theta_prediction = X_test_new.dot(theta_best_values)
```

```
#generate scatterplot
```

```
plt.xlabel("LR predicted")
```

```
plt.ylabel("theta predicted")
```

```
plt.scatter(y_test, theta_prediction)
```

```
plt.show()
```

```
# Compute and print R^2, RMSE, MAE and MSE
```


```
print("R^2: {}".format(lr.score(X_test, y_test)))
```

```
print("Mean Squared Error: {}".format(mean_squared_error(y_test,  
theta_prediction)))
```

```
print("Mean Absolute Error: {}".format(mean_absolute_error(y_test,  
theta_prediction)))
```

```
rmse = np.sqrt(mean_squared_error(y_test, theta_prediction))
```

```
print("Root Mean Squared Error: {}".format(rmse))
```



$R^2$ : 0.9890046246741234  
 Mean Squared Error: 79.81305165097464  
 Mean Absolute Error: 7.228148653430841  
 Root Mean Squared Error: 8.933815066978644

What is the limitation of using the Normal equation for regression?

- Using the normal equation can become very slow for larger data sets and for larger number of features

*#Gradient descent*

cost\_list = []

epoch\_list = []

predicted\_list = []

eta = 0.1 *# learning rate*

n\_iterations = 1000

m = 350

theta = np.random.randn(5,1) *# random initialization*

**for** iteration **in** range(n\_iterations):

    gradients = 2/m \* X\_b.T.dot(X\_b.dot(theta) - y\_new)

    theta = theta - eta \* gradients

    y\_predicted = np.dot(theta.T, X\_b.T)

    cost = np.mean(np.square(y\_new - y\_predicted)) *# MSE (Mean Squared Error)*

```

if iteration%10==0:
    cost_list.append(cost)
    epoch_list.append(iteration)

print("Gradient theta\n" + str(theta))

```

```

Gradient theta
[[5.00531668e+02]
 [2.57625266e+01]
 [3.83285520e+01]
 [1.92209922e-01]
 [6.11735571e+01]]

```


Are they very close to the sklearn's linear regression?

- Yes theta values are extremely close

```

#plot step number against cost
plt.xlabel("epoch")
plt.ylabel("cost")
plt.plot(epoch_list, cost_list)
plt.show()

```



```

#generate prediction based on gradient thetas
gradient_theta_prediction = X_test_new.dot(theta)

```

```

#generate scatterplot
plt.xlabel("LR predicted")
plt.ylabel("theta predicted")

```

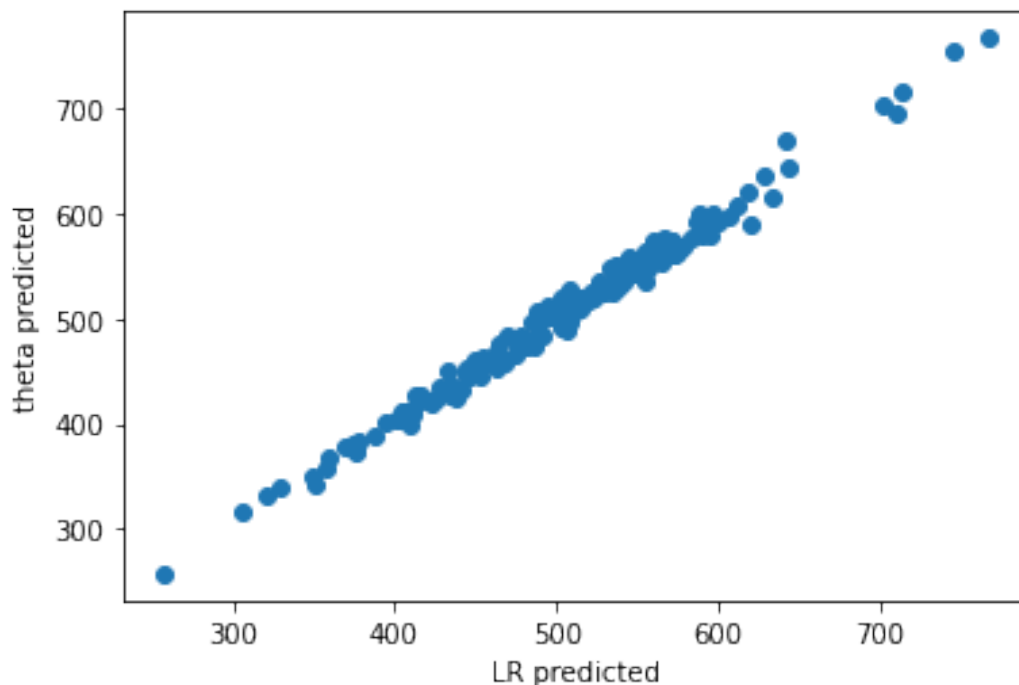
```
plt.scatter(y_test, gradient_theta_prediction)
plt.show()

# Compute and print R^2, RMSE, MAE and MSE
print("R^2: {}".format(lr.score(X_test, y_test)))

print("Mean Squared Error: {}".format(mean_squared_error(y_test,
gradient_theta_prediction)))

print("Mean Absolute Error: {}".format(mean_absolute_error(y_test,
gradient_theta_prediction)))

rmse = np.sqrt(mean_squared_error(y_test, gradient_theta_prediction))
print("Root Mean Squared Error: {}".format(rmse))
```



```
R^2: 0.9890046246741234
Mean Squared Error: 79.81305165097444
Mean Absolute Error: 7.2281486534308295
Root Mean Squared Error: 8.933815066978633
```

Short Question: How do derivatives help in the process of gradient descent?

- If the slope decreases the alpha derivative will also decrease which reduces the step size automatically

Short Question: What are the benefits and the limitations of using batch gradient descent?

- Batch gradient descent is guaranteed to approach the global minimum, or the solution, assuming enough time and a good learning rate

- Batch gradient descent can become very slow with large amounts of data since it uses the whole data set while training, though it scales better with the number of features unlike the normal equation.
- Should preform a hyper parameter optimization to find a good learning rate hyperparameter or the BGD will either take too long to find the solution or constantly over shoot
- If there are multiple features you need to ensure all features have a similar scale or convergence will take much longer
- Irregular data can have local minimums which can slow down the process

*#Stochastic Gradient Descent*

`n_epochs = 50`

`t0, t1 = 5, 50 # learning schedule hyperparameters`

```
def learning_schedule(t):
    return t0 / (t + t1)
```

`theta = np.random.randn(5,1) # random initialization`

```
for epoch in range(n_epochs):
    for i in range(m):
        random_index = np.random.randint(m)
        xi = X_b[random_index:random_index+1]
        yi = y_new[random_index:random_index+1]
        gradients = 2 * xi.T.dot(xi.dot(theta) - yi)
        eta = learning_schedule(epoch * m + i)
        theta = theta - eta * gradients
theta_first = theta
print("Stochastic Gradient theta\n" + str(theta_first))
```

```
Stochastic Gradient theta
[[5.00554695e+02]
 [2.57827596e+01]
 [3.83005549e+01]
 [2.56321105e-02]
 [6.13485039e+01]]
```


Are they very close to the sklearn's linear regression

- Yes theta values are extremely close

*#plot step number against cost*

```
plt.xlabel("epoch")
plt.ylabel("cost")
plt.plot(epoch_list, cost_list)
plt.show()
```





```
#generate prediction based on Stochastic gradient thetas
Stochastic_gradient_theta_prediction = X_test_new.dot(theta_first)


#generate scatterplot
plt.xlabel("LR predicted")
plt.ylabel("theta predicted")
plt.scatter(y_test, Stochastic_gradient_theta_prediction)
plt.show()

# Compute and print R^2, RMSE, MAE and MSE
print("R^2: {}".format(lr.score(X_test, y_test)))

print("Mean Squared Error: {}".format(mean_squared_error(y_test,
Stochastic_gradient_theta_prediction)))

print("Mean Absolute Error: {}".format(mean_absolute_error(y_test,
Stochastic_gradient_theta_prediction)))

rmse = np.sqrt(mean_squared_error(y_test,
Stochastic_gradient_theta_prediction))
print("Root Mean Squared Error: {}".format(rmse))
```



$R^2$ : 0.9890046246741234  
 Mean Squared Error: 79.94079056737529  
 Mean Absolute Error: 7.2326289910661306  
 Root Mean Squared Error: 8.940961389435438

Short Question: What are the benefits and the limitations of using Stochastic gradient descent?

- Possible to train on huge amounts of data sets since it randomly picks only one instance in the training set for each step, much faster and less memory used
- However this also means that it will converge to the minimum via average instead of gently decreasing, once arriving it won't stop making the final values not optimal.
- SGD is actually better for irregular data since it has a better chance of leaving a local minima and instead finding the global minimum

*#SGDRegressor*

```
sgd_reg = SGDRegressor(max_iter=1000, tol=1e-3, penalty=None,
eta0=0.1)
sgd_reg.fit(X_train, y_train)
```

```
print("SGD intercept\n" + str(sgd_reg.intercept_))
print("SGD coeff\n" + str(sgd_reg.coef_))
```

SGD intercept

[499.3792388]

SGD coeff

[26.44472664 38.19148203 -0.50325016 60.99652463]

```

#generate prediction based on gradient thetas
SGDRegressor_theta_prediction = sgd_reg.predict(X_test)

#generate scatterplot
plt.xlabel("LR predicted")
plt.ylabel("theta predicted")
plt.scatter(y_test, SGDRegressor_theta_prediction)
plt.show()


# Compute and print R^2, RMSE, MAE and MSE
print("R^2: {}".format(lr.score(X_test, y_test)))

print("Mean Squared Error: {}".format(mean_squared_error(y_test,
SGDRegressor_theta_prediction)))

print("Mean Absolute Error: {}".format(mean_absolute_error(y_test,
SGDRegressor_theta_prediction)))

rmse = np.sqrt(mean_squared_error(y_test,
SGDRegressor_theta_prediction))
print("Root Mean Squared Error: {}".format(rmse))

```



```

R^2: 0.9890046246741234
Mean Squared Error: 83.2206157055982
Mean Absolute Error: 7.329839552591657
Root Mean Squared Error: 9.122533403917917

```

Briefly explain how mini-batch can overcome the limitations of Batch gradient descent and SGD.

- For each step computing based on small random sets instead of entire set or random instances of one in the set makes it faster than BGD but more accurate than SGD

```
#polynomial of degree 2
```

```
poly_features = PolynomialFeatures(degree=2, include_bias=False)
```

```
X_train_poly = poly_features.fit_transform(X_train)
```

```
X_test_poly = poly_features.fit_transform(X_test)
```

```
#train and predict
```

```
lr.fit(X_train_poly, y_train)
```

```
y_pred = lr.predict(X_test_poly)
```

```
#print coefficients and intercepts
```

```
print("Coefficients\n" + str(lr.coef_))
```

```
print("Intercept " + str(lr.intercept_))
```

```
#generate scatterplot
```

```
plt.xlabel("y_test")
```

```
plt.ylabel("y_pred", rotation = 0)
```

```
plt.scatter(y_test, y_pred)
```

```
plt.show()
```

```
# Compute and print R^2, RMSE, MAE and MSE
```

```
print("R^2: {}".format(lr.score(X_test_poly, y_test)))
```

```
print("Mean Squared Error: {}".format(mean_squared_error(y_test,  
y_pred)))
```

```
print("Mean Absolute Error: {}".format(mean_absolute_error(y_test,  
y_pred)))
```


```
rmse = np.sqrt(mean_squared_error(y_test, y_pred))
```

```
print("Root Mean Squared Error: {}".format(rmse))
```

```
Coefficients
```

```
[ 2.56324041e+01  3.83930621e+01  1.40993709e-01  6.12010238e+01  
 -8.38261559e-01 -1.61454843e-01 -6.96974957e-02  1.53178741e-01  
  4.04112605e-01 -2.24521389e-01  4.19382892e-02  5.97272276e-01  
 -3.66429892e-01 -4.27833581e-02]
```

```
Intercept 500.49414369135536
```



R<sup>2</sup>: 0.9882421824247092  
Mean Squared Error: 85.34745505511978  
Mean Absolute Error: 7.432195703827265  
Root Mean Squared Error: 9.238368636026589

*#polynomial of degree 3*

```
poly_features = PolynomialFeatures(degree=3, include_bias=False)
```

```
X_train_poly = poly_features.fit_transform(X_train)
```

```
X_test_poly = poly_features.fit_transform(X_test)
```

*#train and predict*

```
lr.fit(X_train_poly, y_train)
```

```
y_pred = lr.predict(X_test_poly)
```

*#print coefficients and intercepts*

```
print("Coefficients\n" + str(lr.coef_))
```

```
print("Intercept " + str(lr.intercept_))
```

*#generate scatterplot*

```
plt.xlabel("y_test")
```

```
plt.ylabel("y_pred", rotation = 0)
```

```
plt.scatter(y_test, y_pred)
```

```
plt.show()
```

*# Compute and print R<sup>2</sup>, RMSE, MAE and MSE*

```
print("R2: {}".format(lr.score(X_test_poly, y_test)))
```

```


print("Mean Squared Error: {}".format(mean_squared_error(y_test,
y_pred)))

print("Mean Absolute Error: {}".format(mean_absolute_error(y_test,
y_pred)))

rmse = np.sqrt(mean_squared_error(y_test, y_pred))
print("Root Mean Squared Error: {}".format(rmse))

Coefficients
[25.66180088 37.73817386  1.71438477 61.53147374 -0.9027249  -
0.18506912
  0.13999095  0.62340699  0.22250795 -0.35852515 -0.23855601
0.65188939
 -0.62388113  0.0795928   0.08363636 -0.35101308 -0.10954996
0.40875869
  0.13211318 -0.64021488  0.28780793 -0.28538261  0.66323145  -
0.36183884
 -0.09354415 -0.48873914 -0.28292754  1.1650947   0.55789041
0.47209095
 -0.34630194 -1.07106588 -0.41258377  0.15717302]
Intercept 500.5258833779942

```



```

R^2: 0.986916696696579
Mean Squared Error: 94.96886930852037
Mean Absolute Error: 7.6958010080256205
Root Mean Squared Error: 9.745197243181913

```

```

#learning curve
def plot_learning_curves(model, X, y):
    X_train, X_val, y_train, y_val = train_test_split(X, y, test_size
= .3, random_state = 101)

    train_errors, val_errors = [], []

    for m in range(1, len(X_train) + 1):
        model.fit(X_train[:m], y_train[:m])
        y_train_predict = model.predict(X_train[:m])
        y_val_predict = model.predict(X_val)
        train_errors.append(mean_squared_error(y_train[:m],
y_train_predict))
        val_errors.append(mean_squared_error(y_val, y_val_predict))


    plt.plot(np.sqrt(train_errors), "r-+", linewidth=2, label="train")
    plt.plot(np.sqrt(val_errors), "b-", linewidth=3, label="val")
    plt.legend(loc="upper right", fontsize=14)
    plt.xlabel("Training set size", fontsize=14)
    plt.ylabel("RMSE", fontsize=14)

#plotting linear regression with learning curve
polynomial_regression = Pipeline([
    ("poly_features", PolynomialFeatures(degree=5,
include_bias=False)),
    ("lin_reg", LinearRegression()),
])

plot_learning_curves(polynomial_regression, X, y)

plt.show()


```



```
poly_features = PolynomialFeatures(degree=3, include_bias=False)
```

```
lin_reg = LinearRegression()  
plot_learning_curves(lin_reg, X, y)
```

```
plt.show()
```





Interpret the result

- Function of the learning curves of degree 10 and degree 3 polynomials
- For degree 3 as more instances are added it becomes harder to fit the data perfectly so the train's RMSE increases until it eventually levels out similarly to the val line.

Explain the purpose of regularization

- To constrain a model and make it simpler and reduce the risk of overfitting; less degrees of freedom means harder to overfit data

*#polynomial of degree 3 Regularization*

```
poly_features = PolynomialFeatures(degree=3, include_bias=False)
```

```
X_train_poly = poly_features.fit_transform(X_train)
```

```
X_test_poly = poly_features.fit_transform(X_test)
```

*#ridge Regression*

```
ridge_reg = Ridge(alpha=1, solver="cholesky", random_state=42)
```

*#train and predict*

```
ridge_reg.fit(X_train_poly, y_train)
```

```
y_pred = ridge_reg.predict(X_test_poly)
```

*#generate scatterplot*

```
plt.xlabel("y_test")
```

```
plt.ylabel("y_pred", rotation = 0)
```

```
plt.scatter(y_test, y_pred)
```

```
plt.show()
```

*# Compute and print R^2, RMSE, MAE and MSE*


```
print("R^2: {}".format(lr.score(X_test_poly, y_test)))
```

```
print("Mean Squared Error: {}".format(mean_squared_error(y_test, y_pred)))
```

```
print("Mean Absolute Error: {}".format(mean_absolute_error(y_test, y_pred)))
```

```
rmse = np.sqrt(mean_squared_error(y_test, y_pred))
```

```
print("Root Mean Squared Error: {}".format(rmse))
```



R<sup>2</sup>: 0.986916696696579  
Mean Squared Error: 97.65959147337  
Mean Absolute Error: 7.855335064312333  
Root Mean Squared Error: 9.882286753245424

*#SGDregressor for ridge*

sgd\_reg = SGDRegressor(penalty="l2") *#L2*

*#train and predict*

sgd\_reg.fit(X\_train\_poly, y\_train)

y\_pred = sgd\_reg.predict(X\_test\_poly)

*#generate scatterplot*

plt.xlabel("y\_test")

plt.ylabel("y\_pred", rotation = 0)

plt.scatter(y\_test, y\_pred)

plt.show()


*# Compute and print R<sup>2</sup>, RMSE, MAE and MSE*

print("R<sup>2</sup>: {}".format(lr.score(X\_test\_poly, y\_test)))

print("Mean Squared Error: {}".format(mean\_squared\_error(y\_test, y\_pred)))

print("Mean Absolute Error: {}".format(mean\_absolute\_error(y\_test, y\_pred)))

```
rmse = np.sqrt(mean_squared_error(y_test, y_pred))
print("Root Mean Squared Error: {}".format(rmse))
```



```
R^2: 0.986916696696579
Mean Squared Error: 138.16301767995463
Mean Absolute Error: 9.487205443326394
Root Mean Squared Error: 11.754276569825752
```

```
#Lasso regression
```

```
lasso_reg = Lasso(alpha=0.1)
```

```
#train and predict
```

```
lasso_reg.fit(X_train_poly, y_train)
```

```
y_pred = lasso_reg.predict(X_test_poly)
```

```
#generate scatterplot
```

```
plt.xlabel("y_test")
```

```
plt.ylabel("y_pred", rotation = 0)
```

```
plt.scatter(y_test, y_pred)
```

```
plt.show()
```


```
# Compute and print R^2, RMSE, MAE and MSE
```

```
print("R^2: {}".format(lr.score(X_test_poly, y_test)))
```

```
print("Mean Squared Error: {}".format(mean_squared_error(y_test,
y_pred)))
```

```
print("Mean Absolute Error: {}".format(mean_absolute_error(y_test,
y_pred)))
```

```
rmse = np.sqrt(mean_squared_error(y_test, y_pred))
print("Root Mean Squared Error: {}".format(rmse))
```



```
R^2: 0.986916696696579
Mean Squared Error: 93.21215816943548
Mean Absolute Error: 7.679627716794205
Root Mean Squared Error: 9.654644383375054
```

How Lasso perform the regularization and how does that affect the thetas?

- adds regularization term to cost function using the norm of the weight vector
- automatically eliminates the weights of the least important features by setting them to 0

```
#Elastic net regression
```

```
elastic_net = ElasticNet(alpha=0.1, l1_ratio=0.5, random_state=42)
```

```
#train and predict
```

```
elastic_net.fit(X_train_poly, y_train)
y_pred = elastic_net.predict(X_test_poly)
```

```
#generate scatterplot
```

```
plt.xlabel("y_test")
plt.ylabel("y_pred", rotation = 0)
plt.scatter(y_test, y_pred)
plt.show()
```

```
# Compute and print R^2, RMSE, MAE and MSE
```

```


print("R^2: {}".format(lr.score(X_test_poly, y_test)))

print("Mean Squared Error: {}".format(mean_squared_error(y_test,
y_pred)))

print("Mean Absolute Error: {}".format(mean_absolute_error(y_test,
y_pred)))

rmse = np.sqrt(mean_squared_error(y_test, y_pred))
print("Root Mean Squared Error: {}".format(rmse))

```



```

R^2: 0.986916696696579
Mean Squared Error: 261.69175128224043
Mean Absolute Error: 11.378641625545779
Root Mean Squared Error: 16.17688941923757

```

How ElasticNet different compared to Lasso and RIDGE perform the regularization and how does that affect the thetas?

- Is a mix of both ridge and lasso via a ratio  $r$ , when  $r=0$  then equivalent to ridge, while  $r=1$  then equivalent to lasso. Behaves less erratically when there are more features than training instances or when features are strongly correlated
- Also automatically reduces the weights of useless features