

## This jupyter notebook is prepared by “Nawras Rawas Qalaji”.

#1. Load Data and perform basic EDA

*#I. import neccessary libriaries*

```
import pandas as pd
import numpy as np
import matplotlib as mpl
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
import missingno as msno
import scipy.stats as st
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LinearRegression, SGDRegressor,
Ridge, ElasticNet, Lasso
from sklearn.metrics import classification_report, confusion_matrix
from sklearn.pipeline import Pipeline
from sklearn.cluster import KMeans
from sklearn.metrics.pairwise import euclidean_distances
import scipy.cluster.hierarchy as shc
```

*#II. import hrdata.csv to hrdata dataframe*

```
hrdata = pd.read_csv("/content/drive/MyDrive/Colab
Notebooks/hrdata3.csv")
```

*#count columns and rows*

```
print("Count Columns and rows")
print("columns " + str(hrdata.shape[1]))
print("rows " + str(hrdata.shape[0]) + "\n")
```

*#III. top and bottom 5 rows*

```
print("\nTop Five Rows")
print(hrdata.head())
print("\nBottom Five Rows")
print(hrdata.tail())
```

*#IV. show any missing values*

```
print("\nMissing Values Numerically")
missingValues = hrdata.isnull().sum().sort_values(ascending = False)
print(missingValues)
```

*#V. check data types*

```
print(hrdata.info())
```

Count Columns and rows

columns 8

rows 12977

### Top Five Rows

```

    Unnamed: 0  enrollee_id  city_development_index  experience
company_size  \
0             1         29725                0.776         15
2
1             4          666                0.767         21
2
2             6        28806                0.920          5
2
3             7          402                0.762         13
0
4             8        27107                0.920          7
2

```

```

    last_new_job  training_hours  target
0             5             47      0.0
1             4             8       0.0
2             1            24      0.0
3             5            18      1.0
4             1            46      1.0

```

### Bottom Five Rows

```

    Unnamed: 0  enrollee_id  city_development_index  experience  \
12972        19149         251                0.920          9
12973        19150        32313                0.920         10
12974        19152        29754                0.920          7
12975        19155        24576                0.920         21
12976        19156         5756                0.802          0

```

```

    company_size  last_new_job  training_hours  target
12972          2             1             36      1.0
12973          3             3             23      0.0
12974          1             1             25      0.0
12975          2             4             44      0.0
12976          4             2             97      0.0

```

### Missing Values Numerically

```

    Unnamed: 0      0
enrollee_id      0
city_development_index  0
experience        0
company_size      0
last_new_job      0
training_hours    0
target           0
dtype: int64
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 12977 entries, 0 to 12976

```

Data columns (total 8 columns):

#	Column	Non-Null Count	Dtype
0	Unnamed: 0	12977 non-null	int64
1	enrollee_id	12977 non-null	int64
2	city_development_index	12977 non-null	float64
3	experience	12977 non-null	int64
4	company_size	12977 non-null	int64
5	last_new_job	12977 non-null	int64
6	training_hours	12977 non-null	int64
7	target	12977 non-null	float64

dtypes: float64(2), int64(6)

memory usage: 811.2 KB

None

1) IV. Is there any null values on any column?

- There are no missing or null values

1) V. Are all the columns numeric such as float or int? If not, please convert them to numeric (int/float) before going to the next step.

- Yes all the columns are numeric

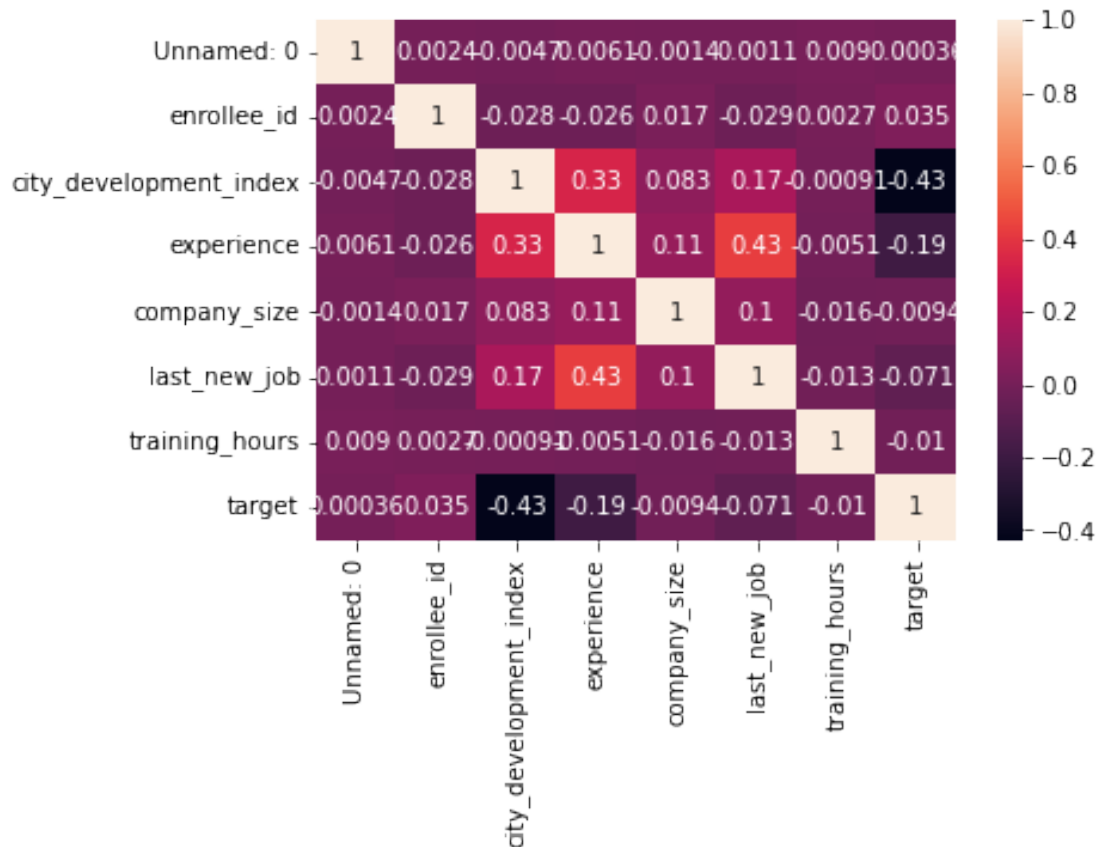
*#VI. Plot the heatmap with correlations*

```
correlationNums = hrdata.select_dtypes(include = [np.number])
```

```
sns.heatmap(data = correlationNums.corr(), annot = True)
```

```
plt.figure()
```

```
plt.show()
```



<Figure size 432x288 with 0 Axes>

## #2. Feature Selection and Pre-processing

*#1. Put all the data from the dataframe into X, except the enrollee\_id and the target columns*

```
hrdata = hrdata.drop(["Unnamed: 0"], axis = 1)
X = hrdata.copy()
X = X.drop(["enrollee_id", "target"], axis = 1)
```

X

	city_development_index	experience	company_size	last_new_job
0	0.776	15	2	5
1	0.767	21	2	4
2	0.920	5	2	1
3	0.762	13	0	5
4	0.920	7	2	1

...	...	...	...	...
12972	0.920	9	2	1
12973	0.920	10	3	3
12974	0.920	7	1	1
12975	0.920	21	2	4
12976	0.802	0	4	2

	training_hours
0	47
1	8
2	24
3	18
4	46
...	...
12972	36
12973	23
12974	25
12975	44
12976	97

[12977 rows x 5 columns]

## *#II. Perform feature scaling*

```
scaler = StandardScaler();
scaler.fit(X)
scaled_X = scaler.transform(X)
reverseScaled_X = scaler.inverse_transform(scaled_X)
```

```
print("\n scaled");
print(scaled_X)
print("\n reverse scaled");
print(reverseScaled_X)
```

scaled
[ [-0.50342203  0.63395707 -0.5747232  1.69076217 -0.30839586]
[ [-0.57841303  1.54600905 -0.5747232  1.08113696 -0.95180478]
[ [ 0.69643399 -0.88612956 -0.5747232 -0.74773864 -0.68784215]
...
[ [ 0.69643399 -0.58211224 -1.0314958 -0.74773864 -0.67134448]
[ [ 0.69643399  1.54600905 -0.5747232  1.08113696 -0.35788885]
[ [-0.28678136 -1.64617288  0.338822 -0.13811344  0.51648738]]

```
reverse scaled
[[ 0.776 15.    2.    5.    47.   ]
 [ 0.767 21.    2.    4.    8.    ]
 [ 0.92   5.    2.    1.    24.   ]
 ...
 [ 0.92   7.    1.    1.    25.   ]
 [ 0.92  21.    2.    4.    44.   ]
 [ 0.802  0.    4.    2.    97.   ]]
```

### #3. KMeans Clustering

*#I. Import related library for Kmeans and perform Kmeans on X*

```
kmeans = KMeans(n_clusters=2, random_state=47, init = "k-means++")
k = kmeans.fit(scaled_X)
```

*#II. Show the cluster centers as it is and then inverse the scale and show the centers*

```
print("\nscaled cluster centers")
print(k.cluster_centers_)
print(k.fit_predict(scaled_X))
```

```
print("\nreversed scaled cluster centers")
```

```
k_reverse = kmeans.fit(reverseScaled_X)
print(k_reverse.cluster_centers_)
print(k_reverse.fit_predict(reverseScaled_X))
```

```
k = kmeans.fit(scaled_X)
```

scaled cluster centers

```
[[-0.31364817 -0.63941844 -0.15207764 -0.55076921  0.01170319]
 [ 0.44177356  0.90062111  0.21420141  0.77575864 -0.01648395]]
[1 1 0 ... 0 1 0]
```

reversed scaled cluster centers

```
[ [ 0.83639888 10.83377685  3.26905531  2.23474625 41.97690553]
 [ 0.83649939 10.81099796  3.21181263  2.19144603 167.34052953]]
[0 0 0 ... 0 0 0]
```

II. Please explain in words about the centers relating them to the columns of the data set

- Kmeans determines the centers by trying to separate data samples into groups of equal variance which is determined by the data in the columns. Essentially similar data samples will be clustered closer together while differing data will be farther apart

*#III. Show the distance matrix*

```
print("\ndistance matrix")
print(k.transform(scaled_X))
```

*#IV. Show the labels*

```
print("\nlabels")
print(k.labels_)
```

*#V. Add a new column to your data frame called cluster\_label*

```
hrdata["cluster_label"] = k.labels_
hrdata
```

distance matrix

```
[[2.6387601  1.58409296]
 [2.93535149  1.74579189]
 [1.3371321   2.57904978]
 ...
 [1.51732572  2.56222519]
 [2.96221279  1.14612076]
 [1.2962834   2.85511871]]
```

labels

```
[1 1 0 ... 0 1 0]
```

	enrollee_id	city_development_index	experience	
company_size \				
0	29725	0.776	15	2
1	666	0.767	21	2
2	28806	0.920	5	2
3	402	0.762	13	0
4	27107	0.920	7	2
...	...	...	...	...
12972	251	0.920	9	2
12973	32313	0.920	10	3
12974	29754	0.920	7	1
12975	24576	0.920	21	2
12976	5756	0.802	0	4

```
last_new_job  training_hours  target  cluster_label
```

0	5	47	0.0	1
1	4	8	0.0	1
2	1	24	0.0	0
3	5	18	1.0	1
4	1	46	1.0	0
...	...	...	...	...
12972	1	36	1.0	0
12973	3	23	0.0	1
12974	1	25	0.0	0
12975	4	44	0.0	1
12976	2	97	0.0	0

[12977 rows x 8 columns]

*#VI. Add a column target\_int and write a function or use a strategy to store the int version of the target column into the target\_int column*

hrdata["target\_int"] = hrdata["target"].astype(int)

hrdata

	enrollee_id	city_development_index	experience	
company_size \				
0	29725	0.776	15	2
1	666	0.767	21	2
2	28806	0.920	5	2
3	402	0.762	13	0
4	27107	0.920	7	2
...	...	...	...	...
12972	251	0.920	9	2
12973	32313	0.920	10	3
12974	29754	0.920	7	1
12975	24576	0.920	21	2
12976	5756	0.802	0	4

	last_new_job	training_hours	target	cluster_label	target_int
0	5	47	0.0	1	0
1	4	8	0.0	1	0



2	1	24	0.0	0	0
3	5	18	1.0	1	1
4	1	46	1.0	0	1
...	...	...	...	...	...
12972	1	36	1.0	0	1
12973	3	23	0.0	1	0
12974	1	25	0.0	0	0
12975	4	44	0.0	1	0
12976	2	97	0.0	0	0

[12977 rows x 9 columns]

*#VII. top and bottom 5 rows with new added columns*

```
print("\nTop Five Rows")
print(hrdata.head())
print("\nBottom Five Rows")
print(hrdata.tail())
```

Top Five Rows

	enrollee_id	city_development_index	experience	company_size	\
0	29725	0.776	15	2	
1	666	0.767	21	2	
2	28806	0.920	5	2	
3	402	0.762	13	0	
4	27107	0.920	7	2	

	last_new_job	training_hours	target	cluster_label	target_int
0	5	47	0.0	1	0
1	4	8	0.0	1	0
2	1	24	0.0	0	0
3	5	18	1.0	1	1
4	1	46	1.0	0	1

Bottom Five Rows

	enrollee_id	city_development_index	experience	company_size	\
12972	251	0.920	9	2	

12973	32313	0.920	10	3
12974	29754	0.920	7	1
12975	24576	0.920	21	2
12976	5756	0.802	0	4

	last_new_job	training_hours	target	cluster_label	target_int
12972	1	36	1.0	0	1
12973	3	23	0.0	1	0
12974	1	25	0.0	0	0
12975	4	44	0.0	1	0
12976	2	97	0.0	0	0

*#VIII. compare the cluster label with the ground truth*

*#Confusion matrix and classification report*

```
print("\nconfusion matrix")
print(confusion_matrix(hrdata["target_int"], k.labels_))
print("\nclassification report")
print(classification_report(hrdata["target_int"], k.labels_))
```

*#Missclassification count*

```
print("\nMissclassification count:")
print(confusion_matrix(hrdata["target_int"], k.labels_)[0][1] +
confusion_matrix(hrdata["target_int"], k.labels_)[1][0])
```

confusion matrix

```
[[5835 4860]
 [1747  535]]
```

classification report

	precision	recall	f1-score	support
0	0.77	0.55	0.64	10695
1	0.10	0.23	0.14	2282
accuracy			0.49	12977
macro avg	0.43	0.39	0.39	12977
weighted avg	0.65	0.49	0.55	12977

Missclassification count:  
6607

IX. Discuss the numbers from 3 Viii and any thoughts on it

- The numbers are not very accurate, since the labels are not related to target, kmeans puts similar data into two groups defined as 0 and 1. However target refers to whether a candidate wants to work for a company after training or not.

*#X. Print the inertia*

```
print(k.inertia_)
```

49643.86379769514

XI. What is the elbow method and what is its purpose of it in the case of KMeans clustering?

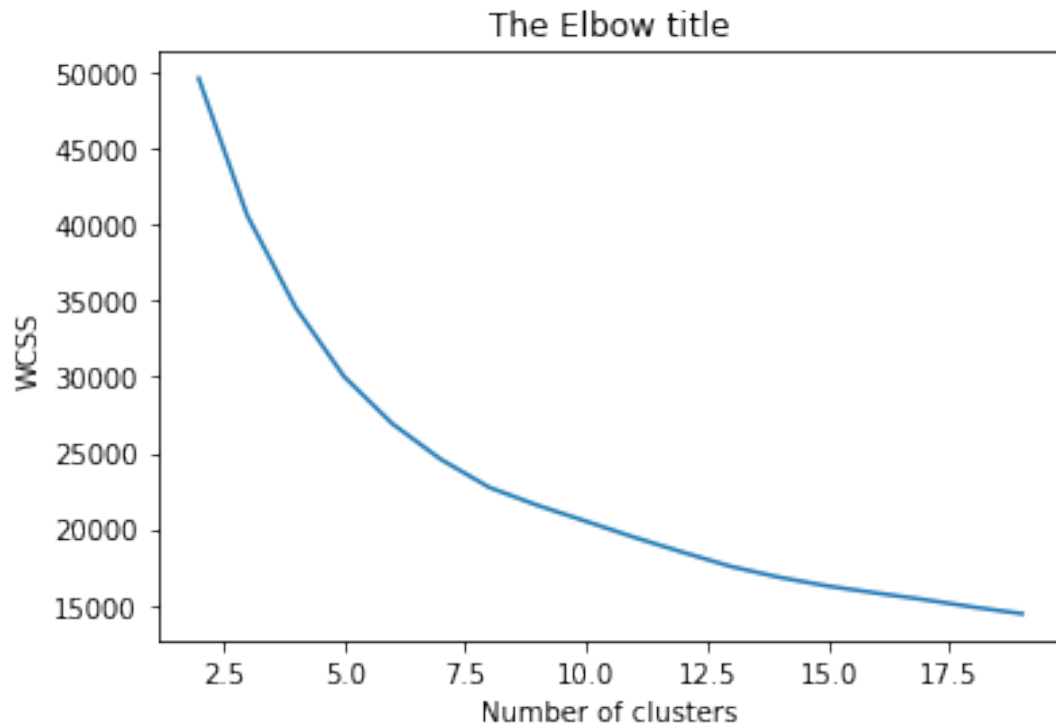
- The elbow method is used to find the optimal number of clusters data should be clustered into. In the case of KMeans it does exactly that, finds the optimal number of k clusters

*#XII Plot the inertia for the different numbers of clusters from 2 to 20*

```
wcss=[]  
for i in range(2,20):  
    kmeans = KMeans(i, init = "k-means++")  
    kmeans.fit(scaled_X)  
    wcss_iter = kmeans.inertia_  
    wcss.append(wcss_iter)
```

```
number_clusters = range(2,20)  
plt.plot(number_clusters,wcss)  
plt.title('The Elbow title')  
plt.xlabel('Number of clusters')  
plt.ylabel('WCSS')
```


```
Text(0, 0.5, 'WCSS')
```



*#XIII Show a scatter plot with training hours against experience where the points should be colored based on the two cluster labels*

```
sns.scatterplot(x = hrdata['training_hours'], y =
hrdata['experience'], hue = hrdata['cluster_label'])
```

<matplotlib.axes.\_subplots.AxesSubplot at 0x7efec5e77650>



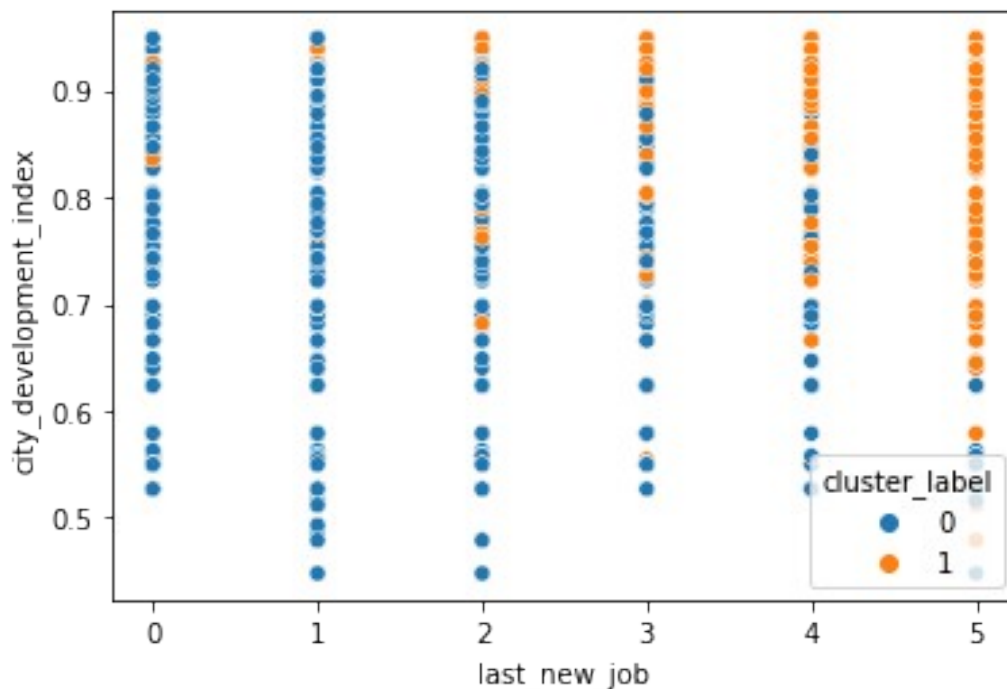
XIII. Write any thoughts on this plot

- The plot shows that data with experience beyond about 10 begin to cluster more towards group 1 and those below are put into group 0. Training hours doesn't appear to have a notable effect on which group the data is separated into compared experience as shown by the plot and how even as training hours increases no notable change is generated.

*#XIII Show a scatter plot with any other two attributes you are interested in*

```
sns.scatterplot(x = hrdata['last_new_job'], y =  
hrdata['city_development_index'], hue = hrdata['cluster_label'])
```

<matplotlib.axes.\_subplots.AxesSubplot at 0x7efec5df8310>



XIV. Write any thoughts on the second plot

- As last new job increases it looks like there is a slight increase in data samples being clustered into group 0, especially towards 5. While there seem to be more instances of a higher city development index the actual number doesn't appear to have an effect on whether it is grouped in group 0 or 1 compared to last new job


#4. AgglomerativeClustering

*#I. plot a dendrogram*

```
plt.figure(figsize= (10,7))
```

```
plt.title("Dendrograms")
```

```
dend = shc.dendrogram(shc.linkage(scaled_X, method='ward'))
```




*#II. Perform AgglomerativeClustering with 2 clusters first, and use euclidean distance for affinity and linkage = 'ward'*

```
from sklearn.cluster import AgglomerativeClustering
AC = AgglomerativeClustering(n_clusters=2, affinity='euclidean',
linkage='ward')
AC.fit_predict(scaled_X)
```

*#III After creating the clusters, plot training hours against experience*

```
sns.scatterplot(x = hrdata['training_hours'], y =
hrdata['experience'], hue = AC.labels_)
```

<matplotlib.axes.\_subplots.AxesSubplot at 0x7efebc3cba50>



### III. Discuss anything interesting

- The data samples grouped in group 1 seem to cluster near the top left portion of the graph, so a higher experience and lower training hours seems to put more data samples into group 1 than anywhere else. This differs from kmeans clustering where experience seemed to have no effect on whether a sample was grouped into group 0 or 1

*#IV. Then, increase the number of clusters to 4 or 5 and build the clusters again and plot them again to see any difference*

```
from sklearn.cluster import AgglomerativeClustering
AC = AgglomerativeClustering(n_clusters=4, affinity='euclidean',
linkage='ward')
AC.fit_predict(scaled_X)
```

```
sns.scatterplot(x = hrdata['training_hours'], y =
hrdata['experience'], hue = AC.labels_)
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7efebc553190>
```



